

*****DESIGN TRAIN AND TEST MULTI LAYER PERCEPTRON FOR TABULAR DATA AND VERIFY VARIOUS ACTIVATION FUNCTIONS AND OPTIMIZERS TENSORFLOW*****

```
# This Python 3 environment comes with many helpful analytics
libraries installed
# It is defined by the kaggle/python Docker image:
https://github.com/kaggle/docker-python
# For example, here's several helpful packages to load

import numpy as np # linear algebra
import pandas as pd # data processing, CSV file I/O (e.g. pd.read_csv)

# Input data files are available in the read-only "../input/"
directory
# For example, running this (by clicking run or pressing Shift+Enter)
will list all files under the input directory

import os
for dirname, _, filenames in os.walk('/kaggle/input'):
    for filename in filenames:
        print(os.path.join(dirname, filename))

# You can write up to 20GB to the current directory (/kaggle/working/)
that gets preserved as output when you create a version using "Save &
Run All"
# You can also write temporary files to /kaggle/temp/, but they won't
be saved outside of the current session

/kaggle/input/diabetes-dataset/diabetes.csv

data=pd.read_csv("/kaggle/input/diabetes-dataset/diabetes.csv")
data
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI
\						
0	6	148	72	35	0	33.6
1	1	85	66	29	0	26.6
2	8	183	64	0	0	23.3

3	1	89	66	23	94	28.1
4	0	137	40	35	168	43.1
..
763	10	101	76	48	180	32.9
764	2	122	70	27	0	36.8
765	5	121	72	23	112	26.2
766	1	126	60	0	0	30.1
767	1	93	70	31	0	30.4

	DiabetesPedigreeFunction	Age	Outcome
0	0.627	50	1
1	0.351	31	0
2	0.672	32	1
3	0.167	21	0
4	2.288	33	1
..
763	0.171	63	0
764	0.340	27	0
765	0.245	30	0
766	0.349	47	1
767	0.315	23	0

[768 rows x 9 columns]

```
import tensorflow as tf
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense
from tensorflow.keras.optimizers import Adam, SGD, RMSprop
from sklearn.datasets import load_diabetes
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import mean_squared_error
import matplotlib.pyplot as plt

# Load and prepare the data
data = load_diabetes()
X = data.data
y = data.target

# Standardize the features
scaler = StandardScaler()
X = scaler.fit_transform(X)
```

```

# Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=42)

# Define the model
def create_model(activation='relu', optimizer='adam'):
    model = Sequential()
    model.add(Dense(64, input_shape=(X_train.shape[1],),
activation=activation))
    model.add(Dense(32, activation=activation))
    model.add(Dense(1)) # No activation for regression output

    # Compile the model
    model.compile(optimizer=optimizer, loss='mean_squared_error',
metrics=['mse'])
    return model

# Define a list of activation functions and optimizers to test
activation_functions = ['relu', 'tanh', 'sigmoid']
optimizers = ['adam', 'sgd', 'rmsprop']

results = {}

# Train and test the model with different configurations
for activation in activation_functions:
    for optimizer in optimizers:
        model = create_model(activation=activation,
optimizer=optimizer)
        model.fit(X_train, y_train, epochs=100, batch_size=32,
verbose=0)

        # Evaluate the model
        loss, mse = model.evaluate(X_test, y_test, verbose=0)
        results[(activation, optimizer)] = mse
        print(f"Activation: {activation}, Optimizer: {optimizer}, MSE:
{mse:.4f}")

# Plotting
plt.figure(figsize=(10, 6))
for activation in activation_functions:
    plt.plot([optimizer for (act, optimizer) in results.keys() if act
== activation],
[results[(activation, optimizer)] for optimizer in
optimizers],
label=f'{activation} activation')

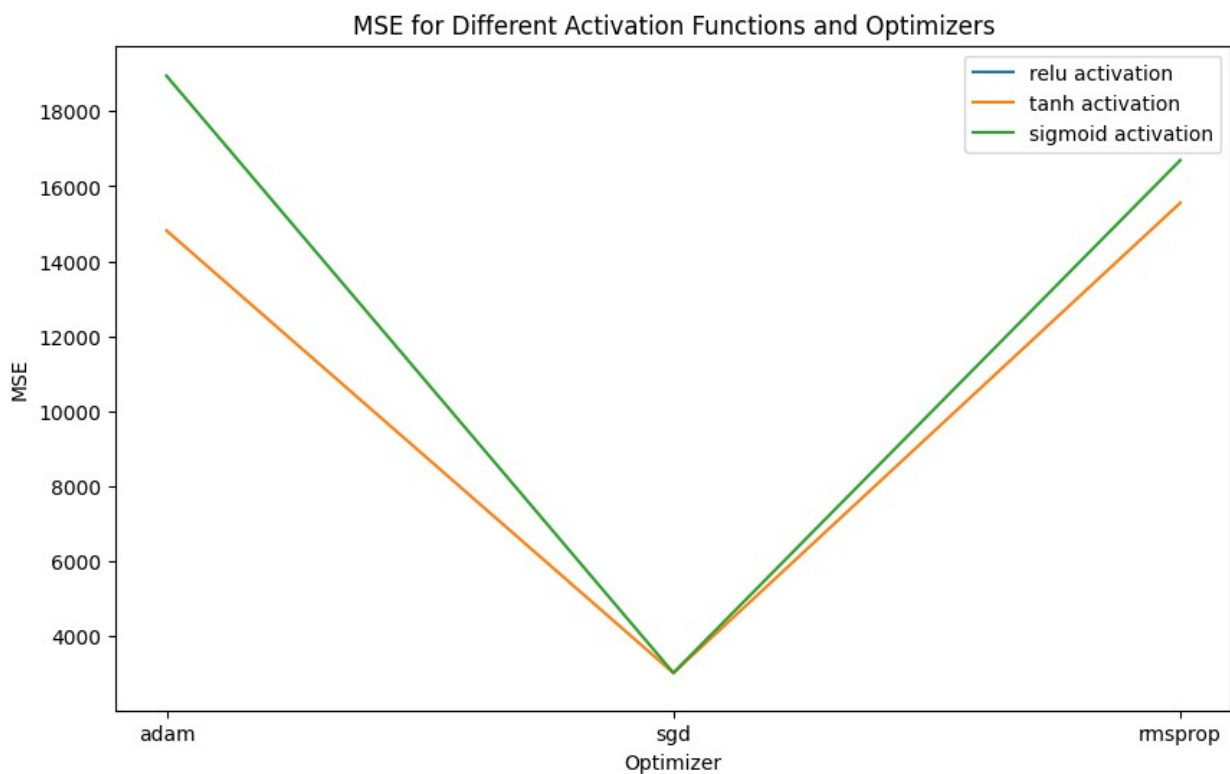
plt.title('MSE for Different Activation Functions and Optimizers')
plt.xlabel('Optimizer')
plt.ylabel('MSE')

```

```
plt.legend()
plt.show()
```

```
/opt/conda/lib/python3.10/site-packages/keras/src/layers/core/
dense.py:87: UserWarning: Do not pass an `input_shape`/`input_dim`
argument to a layer. When using Sequential models, prefer using an
`Input(shape)` object as the first layer in the model instead.
  super().__init__(activity_regularizer=activity_regularizer,
**kwargs)
```

```
Activation: relu, Optimizer: adam, MSE: 2850.3616
Activation: relu, Optimizer: sgd, MSE: nan
Activation: relu, Optimizer: rmsprop, MSE: 2824.1575
Activation: tanh, Optimizer: adam, MSE: 14815.9014
Activation: tanh, Optimizer: sgd, MSE: 3014.3652
Activation: tanh, Optimizer: rmsprop, MSE: 15555.7920
Activation: sigmoid, Optimizer: adam, MSE: 18943.2266
Activation: sigmoid, Optimizer: sgd, MSE: 3018.5166
Activation: sigmoid, Optimizer: rmsprop, MSE: 16693.6094
```



```
from tensorflow.keras.optimizers import SGD
```

```
# Use a smaller learning rate
optimizer = SGD(learning_rate=0.0001)
```

```
# Recreate and compile the model with the adjusted optimizer
```

```

model = create_model(activation='relu', optimizer=optimizer)

# Train the model
model.fit(X_train, y_train, epochs=100, batch_size=32, verbose=0)

# Evaluate the model
loss, mse = model.evaluate(X_test, y_test, verbose=0)
print(f"Activation: relu, Optimizer: sgd, MSE: {mse:.4f}")

Activation: relu, Optimizer: sgd, MSE: 2831.2041

```

To eliminate NaN (Not a Number) values in the MSE when using the ReLU activation function with the SGD optimizer, the following changes can be made:

Reduce Learning Rate: A high learning rate can cause issues with convergence, especially with the ReLU activation function, leading to NaN values. Lowering the learning rate for the SGD optimizer can help stabilize training.

Initialize Weights Properly: Proper weight initialization (like He initialization) can also help prevent issues with ReLU activation.

```

import tensorflow as tf
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense, BatchNormalization
from tensorflow.keras.optimizers import Adam, SGD, RMSprop, Adagrad
from sklearn.datasets import load_diabetes
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import mean_squared_error
import matplotlib.pyplot as plt

# Load and prepare the data
data = load_diabetes()
X = data.data
y = data.target

# Standardize the features
scaler = StandardScaler()
X = scaler.fit_transform(X)

# Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=42)

# Define the model
def create_model(activation='relu', optimizer='adam'):
    model = Sequential()
    model.add(Dense(64, input_shape=(X_train.shape[1],),
activation=activation))
    #model.add(BatchNormalization()) # Add Batch Normalization

```

```

model.add(Dense(32, activation=activation))
#model.add(BatchNormalization()) # Add Batch Normalization
model.add(Dense(1)) # No activation for regression output

# Compile the model with gradient clipping
optimizer = tf.keras.optimizers.get(optimizer)
if isinstance(optimizer, SGD):
    optimizer.learning_rate = 0.0001 # Reduced learning rate

model.compile(optimizer=optimizer, loss='mean_squared_error',
metrics=['mse'])
return model

# Define a list of activation functions and optimizers to test
activation_functions = ['relu', 'tanh', 'sigmoid']
optimizers = ['adam', 'sgd', 'rmsprop', 'Adagrad']

results = {}

# Train and test the model with different configurations
for activation in activation_functions:
    for optimizer in optimizers:
        model = create_model(activation=activation,
optimizer=optimizer)
        model.fit(X_train, y_train, epochs=100, batch_size=32,
verbose=0)

        # Evaluate the model
        loss, mse = model.evaluate(X_test, y_test, verbose=0)
        results[(activation, optimizer)] = mse
        print(f"Activation: {activation}, Optimizer: {optimizer}, MSE:
{mse:.4f}")

# Plotting
plt.figure(figsize=(10, 6))
for activation in activation_functions:
    plt.plot([optimizer for (act, optimizer) in results.keys() if act
== activation],
[results[(activation, optimizer)] for optimizer in
optimizers],
label=f'{activation} activation')

plt.title('MSE for Different Activation Functions and Optimizers')
plt.xlabel('Optimizer')
plt.ylabel('MSE')
# Option 1: Manually set the y-axis limits
#plt.ylim(0, 30000) # Set lower and upper bounds for the y-axis

# Option 2: Use a logarithmic scale for the y-axis
#plt.yscale('log')

```

```
plt.legend()  
plt.show()
```

```
/opt/conda/lib/python3.10/site-packages/keras/src/layers/core/  
dense.py:87: UserWarning: Do not pass an `input_shape`/`input_dim`  
argument to a layer. When using Sequential models, prefer using an  
`Input(shape)` object as the first layer in the model instead.  
    super().__init__(activity_regularizer=activity_regularizer,  
    **kwargs)
```

```
Activation: relu, Optimizer: adam, MSE: 2956.9028  
Activation: relu, Optimizer: sgd, MSE: 3147.7812  
Activation: relu, Optimizer: rmsprop, MSE: 2907.7595  
Activation: relu, Optimizer: Adagrad, MSE: 23059.0273  
Activation: tanh, Optimizer: adam, MSE: 14791.1758  
Activation: tanh, Optimizer: sgd, MSE: 2860.1355  
Activation: tanh, Optimizer: rmsprop, MSE: 15218.6572  
Activation: tanh, Optimizer: Adagrad, MSE: 25984.4414  
Activation: sigmoid, Optimizer: adam, MSE: 20165.0508  
Activation: sigmoid, Optimizer: sgd, MSE: 2957.6797  
Activation: sigmoid, Optimizer: rmsprop, MSE: 16949.9629  
Activation: sigmoid, Optimizer: Adagrad, MSE: 25664.1484
```

